

ECE 420 Parallel and Distributed Programming

Lab 1: Matrix Multiplication with Pthreads

Winter 2017

In this lab, we will implement a parallel program for matrix multiplication with Pthreads. This lab manual mainly consists of two parts: Section 1 and Section 2 introduce the background of the problem and the requirement for this lab, while Section 3 and Section 4 provide some help regarding the working environment and debugging.

1 Background: Matrix Multiplication

Matrix multiplication is a binary operation on two matrices producing a third matrix. Specifically, for an $m \times n$ matrix $A = (a_{ij})$ and an $n \times k$ matrix $B = (b_{ij})$,¹the product $A \cdot B$ is an $m \times k$ matrix $C = (c_{ij})$, where for $\forall i$ and j ,

$$c_{ij} = \sum_{r=0}^{n-1} a_{ir} \cdot b_{rj}. \quad (1)$$

It is natural to think about speeding up the calculation by simultaneously processing several elements in C by several threads. We can make every thread in charge of calculating a certain group of c_{ij} . To balance the load, we can assign the same number of elements in C to each thread (or about the same number if the number of rows/columns is not a multiple of the number of threads). Specifically, we can divide C into blocks of similar sizes and assign the elements in each block (or a submatrix) to a thread, as shown in Fig. 1.

¹All our indices start from 0 to respect the C convention.

0 P ₀₀	1 P ₀₁	2 P ₀₂	3 P ₀₃	4 P ₀₄
5 P ₁₀	6 P ₁₁	7 P ₁₂	8 P ₁₃	9 P ₁₄
10 P ₂₀	11 P ₂₁	12 P ₂₂	13 P ₂₃	14 P ₂₄
15 P ₃₀	16 P ₃₁	17 P ₃₂	18 P ₃₃	19 P ₃₄
20 P ₄₀	21 P ₄₁	22 P ₄₂	23 P ₄₃	24 P ₄₄

Figure 1: Partitioning a matrix into blocks

To simplify the implementation, we only consider the square matrices, i.e., all A , B and C are $n \times n$ matrices. Furthermore, we only consider the number of threads p by which n^2 is **divisible** and we only consider the **square** numbers for p . Then, the matrix C is divided into an array of $\sqrt{p} \times \sqrt{p}$ blocks. Denote the block in the x^{th} row and the y^{th} column as P_{xy} . For a thread of rank k , $0 \leq k \leq p - 1$, we can map it to the block P_{xy} , where $x = \lfloor \frac{k}{\sqrt{p}} \rfloor$ and $y = k \% \sqrt{p}$.² Moreover, the thread of rank k mapped to block P_{xy} contains the elements c_{ij} for $\forall i, j$ such that

$$\frac{n}{\sqrt{p}}x \leq i \leq \frac{n}{\sqrt{p}}(x + 1) - 1$$

and

$$\frac{n}{\sqrt{p}}y \leq j \leq \frac{n}{\sqrt{p}}(y + 1) - 1.$$

² “ $\lfloor \cdot \rfloor$ ” returns the floor; “ $\%$ ” is modulus (returns the remainder after division).

2 Tasks and Requirements

Task:

Implement a shared-memory parallel program for matrix multiplication using Pthreads with the block partition strategy described in Section 1.

Requirements:

1. Use the scripts in “Development Kit Lab 1” to generate input data, load data and save result. Refer to the *readme* file for details on how to use them.
2. Time measurement should be implemented in your code and saved through the provided IO functions. Refer to Section 3.3 for details.
3. The number of threads should be passed as the only command line argument to your program.

Submission:

Each team is required to submit a zip file to eClass by the corresponding due date. The zip file should be named “StudentID-Hx.zip”, where “StudentID” is the Student ID of **one** of your group members (doesn’t matter which member) and “Hx” is the section (H1 or H2).

The zip file should contain the following files:

1. “readme”: a text file containing instructions on how to compile your source files;
2. “members”: a text file listing the student IDs of ALL group members, with each student ID occupying one line;
3. “main”: your final *executable* program file;
4. All the necessary source files to build the executable “main”;

DO NOT include the compiled data generation file, “serialtester” file, or the input/output data file.

Note: you **MUST** use the file names suggested above. File names are case-sensitive. You **MUST** generate the required zip file by directly

compressing all the above files, rather than compressing a folder containing those files.

Example:

Suppose Jack with student ID 1234567 and Rose with student ID 7654321 are in a team in Session H1. They should submit a zip file named “1234567-H1.zip”, in which it contains “readme”, “members”, “main” and all the necessary source files to build “main”. The “members” file should be a file with two lines, the first line being “1234567” and the second line being “7654321”.

Hints:

The correctness of your results can be checked by the script “serialtester.c” in the “Development Kit Lab 1”

3 Basics on Compiling and Running Programs on Linux Systems

3.1 Editing, Compiling and Executing the Code

Pthreads is not a programming language, but an extension package. In this lab, we will use C. You can use whatever text editor you like to write the code. Some simple text editor like VIM, emacs and gedit will be good enough. To use the Pthreads package, you need to include the header file “pthread.h” in your code:

```
#include <pthread.h>
```

After you finish your code, you need to compile it to generate the executable file. It is more or less the same as compiling a typical C program in Linux. For example, if the code file name is “demo.c”, in the terminal, supposing the current path is your code folder, the command to compile will be

```
$gcc -g -Wall -o demo demo.c -lpthread
```

“-g” will generate the necessary information for debugger. “-Wall” will turn on all the warnings. “-o demo” specifies the output file path and name. “demo.c” is the source code file. “-lpthread” tells the compiler to link the Pthreads library. Once it is successfully compiled and error free. You can execute the code by

```
$/demo <possible command line parameters>
```

3.2 Starting and Terminating a Thread

The Pthreads library can create thread to run a function. The function has a special prototypes for the Pthreads. It has a “void*” return type and the argument is a void pointer “void*”. For example, supposing the name of the function is “threadfunc”, the prototype should be

```
void* threadfunc (void* arg_p)
```

In fact we can pass whatever argument through this pointer “arg_p”. In this lab, we only focus on the single program multiple data scheme, so typically we will pass the rank into the thread function. Inside the thread function, we need to cast the void type pointer back to the desired type before we can get access to those arguments.

Pthreads uses “pthread_t” data structure to store the thread information and handle them. We need to assign each thread an individual “pthread_t” object. Same as all the other variables in C language, we need to declare the “pthread_t” objects before we use them.

To start a thread running a specific function, we use pthread_create. The syntax is

```
int pthread_create (  
pthread_t* thread_p  
const pthread_attr_t* attr_p  
void* (*start_routine)(void*)  
void* arg_p)
```

“thread_p” is the pointer of the handle we assign to the thread. We don’t use the second attribute in the lab. The third one is the function we want the thread to run. The “arg_p” is the pointer to the argument we want to pass to the thread function. Say if we want to start a thread running the function “threadfunc” with an assigned rank “1”, the following code will do so

```
pthread_t thread_handle; /*Declare the object before you use  
it*/  
int thread_idx=1; /*Here we want to pass the rank ``1`` to the  
thread*/
```

```
/*...some other code*/  
pthread_create(&thread_handle, NULL,  
threadfunc, (void*) thread_idx);
```

Note that we need to cast the type “int” into “void*” for the desired argument we want to pass to the thread function.

We use `pthread_join` to wait the thread to stop in our program and collect the returned arguments by the threads. The syntax is

```
pthread_join(pthread_t thread_p  
void** ret_val_p)
```

In our example, if we simply want to wait for the thread function “threadfunc” with the handle “thread_handle” ignoring the returning value, we can use `pthread_join(thread_handle, NULL);`

3.3 Time Measurement

The primitive motivation of utilizing the parallel approach is to speed up our program. To find out the real performance and for some evaluation purposes, we need to measure the time consumed by the program. Different to the serial program, it will make no sense to use the `clock` function in C since we are more interested in the total elapsed time not the CPU time.³ It is not suitable to use the linux shell command `time` since it will include all the time the program will take in which some portion like the time consumed by the IO is not of our interest. In the parallel program with Pthreads, we can use some function in the `time.h` header to take down the time at both the begin and end of the main calculating segment. Then we can get the time by the difference of those two time.

The header “timer.h” has defined a macro for time measurement. Refer to the *readme* and the notes in that file for more detail.

In Lab 1, to measure the time in our program, it is enough to record the start time right before you create the threads and record the end time right after you stop the threads by `pthread_join`.

³Actually, we can expect that the parallel program will take more CPU time than the serial one. The parallel version cannot shrink the necessary calculation. To the contrary, it will introduce some overhead and other cost which the serial program will not contain.

4 Debugging and Testing

4.1 Debugging a Parallel Program

Debugging might be one of the toughest part in the parallel programming. There should be thousands of more words than here to go through all the aspects of debugging.

However, never be afraid of the bugs! Be confident in yourself and we can fix everything if we carefully check it. The worst case is only that we check the program line by line with some debugging tools.

On the other hand, although we should be able to fix everything by debugging, it is always better to be more careful in the developing stage and try to prevent the mistakes by good designing.

As for the debugging for parallel program, the challenge is that the thread runs simultaneously and the results are nondeterministic. We can hardly test all the possible situations. Unlike debugging the serial program, debugging a parallel program is state of the art.

One possible approach would be

1. Write your code so that it can run in serial: perfect that first.
2. Deal with communication, synchronization and deadlock on smaller number of threads.
3. Only then try the full size.

Note that in our Pthreads program, assigning only one thread to run the program would be an efficient way to run it in a serial manner. Also, you don't have to follow this approach. You can come up with your own better strategies and you are always welcome to share your ideas.

4.2 Basics on Using `gdb` Debugger

In our lab, you can use whatever debugging tool as you like. Here we will introduce the basics on using the `gdb` debugger. `gdb` is a command line based simple but powerful debugger.

You can launch the `gdb` by the “`gdb`” command in the terminal. If you want to debug the executable program “`demo`”, you can type

```
$ gdb demo -tui
```

Table 1: Basic gdb Commands

Commands	Usage	Example
<code>b</code>	set the break point	<code>b main; b 41</code>
<code>info b</code>	list break points	
<code>delete</code>	delete break point	<code>delete 2</code>
<code>run</code>	run the program	<code>run [args]</code>
<code>n</code>	step to the next statement	
<code>s</code>	step into the function	
<code>c</code>	continue running	
<code>p</code>	display variable value	<code>p V</code>
<code>set</code>	set variable value	<code>set V=3</code>

Note that you need to compile the code with the “-g” flag to link the executable code to the source to use the debugger. The “-tui” option will launch a simple GUI.

Table 4.2 shows the typical commands for debugging a serial program. Note that to set the break points, you can either indicate the function name or the line number of the code. “run” will start running the program and possible command line arguments can be set after it. The index for the breakpoints are generated by the gdb, you can use “info b” to check the index and delete the corresponding breakpoint with “delete”.

As for debugging a program with multiple threads, Table 4.2 shows the basic commands. When you want to debug a multiple thread program, you need to set break points inside the thread function first. The program will stop at the breakpoint in one of the thread. You can then check the current running threads with “info thread”. There will be a “*” before the active thread. You can also find the indexes and you can switch to other threads by the command “thread [thread index]”.

Note that when you are inside a thread, command like “n”, “s” etc. will only influence the current thread and all the other threads will do nothing. However, when you use “c”, all the threads will run simultaneously, and it will stop at a break point of which thread first hits its next breakpoint. This means for example, when you are in Thread 2, after you input the command “c” it could not be in

Table 2: gdb Thread Commands

Commands	Usage	Example
<code>info thread</code>	show the running threads	
<code>thread</code>	switch into another thread	<code>thread 2</code>

Thread 2 when it stops again. Similar situation will occur for “run”, and it will stop in the thread who first hit its breakpoint, rather than Thread 1.

For more commands and information, you can type “help” for more details or check it out online.

4.3 Testing Your Program

Testing is always an important procedure of coding and even tougher than debugging, especially for parallel programming. It is generally considered to be infeasible to thoroughly test a program of a moderate complexity. However, to be a good programmer, we need to try our best to ensure the quality of the program. In our lab, we only have a minimum requirement on testing, i.e., to guarantee the correctness of your program.

Since we cannot cover all the possible inputs, we need to carefully choose the testing cases and justify the correctness as much as possible. Due to the nondeterministic property of parallel programs, some potential errors might not appear at first in some cases. We might need to test a same case for several times for some kinds of parallel programs.

A Appendix: Marking Guideline

Correct submission:	1
Correct results under different thread number settings:	3
Time measurement:	1
Lab Report:	0 (Not required for this lab)
Total:	5